

# ML\_for\_cybersec\_Lab2\_mm11070

December 9, 2022

## 1 Machine Learning for Cybersecurity - Lab 02

Name: Mukta Maheshwari NETID: mm11070

---

### 1.1 Import package

```
[26]: #importing libraries
import pandas as pd
import seaborn as sns
import keras
from tqdm import tqdm
import os
import tarfile
import requests
import re
import sys
import warnings
warnings.filterwarnings('ignore')
import h5py
import numpy as np
import tensorflow as tf
from tensorflow import keras
from keras import backend as K
from keras.models import Model
import matplotlib.pyplot as plt
from mpl_toolkits.axes_grid1.inset_locator import inset_axes
import matplotlib.font_manager as font_manager
import cv2
```

```
[27]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[28]: warnings.filterwarnings("ignore")
```

## 1.2 BadNets

It shows the original badnet and it will print out the accuracy and attack success rate for the original badnet

```
[29]: def data_loader(filepath):
      data = h5py.File(filepath, 'r')
      x_data = np.array(data['data'])
      y_data = np.array(data['label'])
      x_data = x_data.transpose((0,2,3,1))

      return x_data, y_data

[32]: clean_data_filename = '/content/drive/MyDrive/Lab2_mm11070/lab3/data/cl/valid.
      ↪h5'
      poisoned_data_filename = '/content/drive/MyDrive/Lab2_mm11070/lab3/data/bd/
      ↪bd_valid.h5'
      model_filename = '/content/drive/MyDrive/Lab2_mm11070/lab3/models/bd_net.h5'

[33]: #loading data
      def main():
          cl_x_test, cl_y_test = data_loader(clean_data_filename)
          bd_x_test, bd_y_test = data_loader(poisoned_data_filename)

          bd_model = keras.models.load_model(model_filename)

          cl_label_p = np.argmax(bd_model.predict(cl_x_test), axis=1)
          clean_accuracy = np.mean(np.equal(cl_label_p, cl_y_test))*100
          print('Clean Classification accuracy:', clean_accuracy)

          bd_label_p = np.argmax(bd_model.predict(bd_x_test), axis=1)
          asr = np.mean(np.equal(bd_label_p, bd_y_test))*100
          print('Attack Success Rate:', asr)

      if __name__ == '__main__':
          main()
```

```
361/361 [=====] - 8s 20ms/step
Clean Classification accuracy: 98.64899974019225
361/361 [=====] - 7s 19ms/step
Attack Success Rate: 100.0
```

### 1.2.1 Seeing the model structure

```
[34]: model = keras.models.load_model(model_filename)

[35]: print(model.summary())
```

```
Model: "model_1"
```

Layer (type)	Output Shape	Param #	Connected to
input (InputLayer)	[(None, 55, 47, 3)]	0	[]
conv_1 (Conv2D)	(None, 52, 44, 20)	980	['input[0][0]']
pool_1 (MaxPooling2D) ['conv_1[0][0]']	(None, 26, 22, 20)	0	
conv_2 (Conv2D) ['pool_1[0][0]']	(None, 24, 20, 40)	7240	
pool_2 (MaxPooling2D) ['conv_2[0][0]']	(None, 12, 10, 40)	0	
conv_3 (Conv2D) ['pool_2[0][0]']	(None, 10, 8, 60)	21660	
pool_3 (MaxPooling2D) ['conv_3[0][0]']	(None, 5, 4, 60)	0	
conv_4 (Conv2D) ['pool_3[0][0]']	(None, 4, 3, 80)	19280	
flatten_1 (Flatten) ['pool_3[0][0]']	(None, 1200)	0	
flatten_2 (Flatten) ['conv_4[0][0]']	(None, 960)	0	
fc_1 (Dense) ['flatten_1[0][0]']	(None, 160)	192160	
fc_2 (Dense) ['flatten_2[0][0]']	(None, 160)	153760	
add_1 (Add)	(None, 160)	0	['fc_1[0][0]', 'fc_2[0][0]']
activation_1 (Activation)	(None, 160)	0	['add_1[0][0]']
output (Dense) ['activation_1[0][0]']	(None, 1283)	206563	

=====

Total params: 601,643

Trainable params: 601,643

Non-trainable params: 0

-----


-----

None

Visualizing the data to see the clean data

```
[36]: x_data, y_data = data_loader(clean_data_filename)
```


```
[37]: figure = plt.figure(figsize=(10,8))
      cols, rows = 3,3
      for i in range(1, cols*rows+1):
          index = np.random.randint(x_data.shape[0], size=1)
          img, label = (x_data[index], y_data[index])
          figure.add_subplot(rows, cols, i)
          plt.title("true label: {}".format(label))
          plt.axis("off")
          plt.imshow(img[0]/255)
      plt.show()
```



Visualizing the sunglasses poisoned test data

```
[38]: x_poisoned_data, y_poisoned_data = data_loader(poisoned_data_filename)
```

```
[39]: figure = plt.figure(figsize=(10,8))
      cols, rows = 3,3
      for i in range(1, cols*rows+1):
          index = np.random.randint(x_poisoned_data.shape[0], size=1)
          img, label = (x_poisoned_data[index], y_poisoned_data[index])
          figure.add_subplot(rows, cols, i)
          plt.title("true label: {}".format(label))
          plt.axis("off")
          plt.imshow(img[0]/255)
      plt.show()
```



```
[40]: # clearing the session
keras.backend.clear_session()
```

### 1.3 Prune defense

For Pruning the model, the steps are as follows –

1. It is determined whether the final pooling layer, “(pool 3),” is activated.
2. Pruning ALWAYS begins with the smallest average activation. There are 60 total channels in the convolution layer “(conv 3),” and we need to find the index to prune.

```
[41]: # getting the data
cl_x_test, cl_y_test = data_loader(clean_data_filename)
bd_x_test, bd_y_test = data_loader(poisoned_data_filename)
```

```

[42]: clean_data_acc = 98.64899974019225          #from the main function cell - clean_u
      ↪data accuracy
model_copy = keras.models.clone_model(model)
model_copy.set_weights(model.get_weights())
prune_index = []
clean_acc = []
asrate = []
saved_model = np.zeros(3,dtype=bool)

# getting the activation from the last pooling layer
layer_output=model_copy.get_layer('pool_3').output
intermediate_model=keras.models.Model(inputs=model_copy.
      ↪input,outputs=layer_output)
intermediate_prediction=intermediate_model.predict(cl_x_test)
temp = np.mean(intermediate_prediction,axis=(0,1,2))
seq = np.argsort(temp)
weight_0 = model_copy.layers[5].get_weights()[0]
bias_0 = model_copy.layers[5].get_weights()[1]

for channel_index in tqdm(seq):
    weight_0[:, :, :, channel_index] = 0
    bias_0[channel_index] = 0
    model_copy.layers[5].set_weights([weight_0, bias_0])
    cl_label_p = np.argmax(model_copy.predict(cl_x_test), axis=1)
    clean_accuracy = np.mean(np.equal(cl_label_p, cl_y_test))*100
    if (clean_data_acc-clean_accuracy >= 2 and not saved_model[0]):
        print("The accuracy drops at least 2%, saved the model")
        model_copy.save('model_X=2.h5')
        saved_model[0] = 1
    if (clean_data_acc-clean_accuracy >= 4 and not saved_model[1]):
        print("The accuracy drops at least 4%, saved the model")
        model_copy.save('model_X=4.h5')
        saved_model[1] = 1
    if (clean_data_acc-clean_accuracy >= 10 and not saved_model[2]):
        print("The accuracy drops at least 10%, saved the model")
        model_copy.save('model_X=10.h5')
        saved_model[2] = 1
    clean_acc.append(clean_accuracy)
    bd_label_p = np.argmax(model_copy.predict(bd_x_test), axis=1)
    asr = np.mean(np.equal(bd_label_p, bd_y_test))*100
    asrate.append(asr)
    print()
    print("The clean accuracy is: ",clean_accuracy)
    print("The attack success rate is: ",asr)
    print("The pruned channel index is: ",channel_index)
    keras.backend.clear_session()

```

361/361 [=====] - 7s 18ms/step

0%| | 0/60 [00:00<?, ?it/s]  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 7s 20ms/step

2%| | 1/60 [00:15<15:43, 15.99s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 0  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step

3%| | 2/60 [00:31<15:25, 15.96s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 26  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 27

5%| | 3/60 [00:50<16:20, 17.20s/it]  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 8s 21ms/step

7%| | 4/60 [01:06<15:39, 16.77s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 30  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 8s 22ms/step

8%| | 5/60 [01:27<16:49, 18.36s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 31  
361/361 [=====] - 11s 29ms/step  
361/361 [=====] - 8s 21ms/step

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 33



10%| | 6/60 [01:47<16:50, 18.72s/it]  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 9s 24ms/step

12%| | 7/60 [02:05<16:29, 18.67s/it]  
  
The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 34  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 7s 20ms/step

13%| | 8/60 [02:24<16:08, 18.63s/it]  
  
The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 36  
361/361 [=====] - 9s 25ms/step  
361/361 [=====] - 7s 20ms/step

15%| | 9/60 [02:42<15:48, 18.60s/it]  
  
The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 37  
361/361 [=====] - 8s 22ms/step  
361/361 [=====] - 8s 22ms/step

17%| | 10/60 [02:59<15:05, 18.11s/it]  
  
The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 38  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 8s 22ms/step

18%| | 11/60 [03:16<14:28, 17.73s/it]  
  
The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 25  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step

20%| | 12/60 [03:32<13:38, 17.05s/it]  
  
The clean accuracy is: 98.64899974019225

The attack success rate is: 100.0  
The pruned channel index is: 39  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 7s 21ms/step  
22%| | 13/60 [03:50<13:40, 17.47s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 41  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 8s 23ms/step  
23%| | 14/60 [04:10<13:52, 18.10s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 44  
361/361 [=====] - 7s 21ms/step  
361/361 [=====] - 7s 21ms/step  
25%| | 15/60 [04:26<13:04, 17.43s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 45  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step  
27%| | 16/60 [04:41<12:20, 16.82s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 47  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step  
28%| | 17/60 [04:59<12:22, 17.26s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 48  
361/361 [=====] - 10s 27ms/step  
361/361 [=====] - 8s 21ms/step  
30%| | 18/60 [05:20<12:49, 18.33s/it]

The clean accuracy is: 98.64899974019225

The attack success rate is: 100.0  
The pruned channel index is: 49  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step  
32%| | 19/60 [05:38<12:31, 18.32s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 50  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 7s 21ms/step  
33%| | 20/60 [05:54<11:44, 17.61s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 53  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step  
35%| | 21/60 [06:13<11:38, 17.91s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 55  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 8s 21ms/step  
37%| | 22/60 [06:29<10:55, 17.24s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 40  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 7s 19ms/step  
38%| | 23/60 [06:47<10:46, 17.47s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 24  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step  
40%| | 24/60 [07:05<10:39, 17.75s/it]

The clean accuracy is: 98.64899974019225

The attack success rate is: 100.0  
The pruned channel index is: 59  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step  
42%| | 25/60 [07:20<09:54, 16.98s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 9  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 7s 19ms/step  
43%| | 26/60 [07:38<09:47, 17.27s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 2  
361/361 [=====] - 9s 25ms/step  
361/361 [=====] - 7s 20ms/step  
45%| | 27/60 [07:55<09:28, 17.23s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 12  
361/361 [=====] - 8s 22ms/step  
361/361 [=====] - 7s 20ms/step  
47%| | 28/60 [08:11<08:58, 16.84s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 13  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 9s 24ms/step  
48%| | 29/60 [08:31<09:09, 17.72s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 17  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step  
50%| | 30/60 [08:47<08:37, 17.26s/it]

The clean accuracy is: 98.64899974019225

The attack success rate is: 100.0  
The pruned channel index is: 14  
361/361 [=====] - 8s 22ms/step  
361/361 [=====] - 7s 19ms/step  
52%| | 31/60 [09:03<08:07, 16.80s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 15  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 19ms/step  
53%| | 32/60 [09:21<08:01, 17.20s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 23  
361/361 [=====] - 8s 21ms/step  
361/361 [=====] - 7s 19ms/step  
55%| | 33/60 [09:37<07:30, 16.68s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 6  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 10s 26ms/step  
57%| | 34/60 [09:57<07:45, 17.89s/it]

The clean accuracy is: 98.64033948211657  
The attack success rate is: 100.0  
The pruned channel index is: 51  
361/361 [=====] - 8s 22ms/step  
361/361 [=====] - 7s 19ms/step  
58%| | 35/60 [10:13<07:13, 17.34s/it]

The clean accuracy is: 98.64033948211657  
The attack success rate is: 100.0  
The pruned channel index is: 32  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 7s 19ms/step  
60%| | 36/60 [10:32<07:03, 17.66s/it]

The clean accuracy is: 98.63167922404088

The attack success rate is: 100.0  
The pruned channel index is: 22  
361/361 [=====] - 7s 18ms/step  
361/361 [=====] - 7s 18ms/step  
62%| | 37/60 [10:47<06:29, 16.92s/it]

The clean accuracy is: 98.65765999826795  
The attack success rate is: 100.0  
The pruned channel index is: 21  
361/361 [=====] - 7s 18ms/step  
361/361 [=====] - 7s 18ms/step  
63%| | 38/60 [11:01<05:53, 16.07s/it]

The clean accuracy is: 98.64899974019225  
The attack success rate is: 100.0  
The pruned channel index is: 20  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 7s 18ms/step  
65%| | 39/60 [11:15<05:25, 15.52s/it]

The clean accuracy is: 98.6056984498138  
The attack success rate is: 100.0  
The pruned channel index is: 19  
361/361 [=====] - 7s 18ms/step  
361/361 [=====] - 7s 19ms/step  
67%| | 40/60 [11:30<05:03, 15.16s/it]

The clean accuracy is: 98.57105741751104  
The attack success rate is: 100.0  
The pruned channel index is: 43  
361/361 [=====] - 9s 26ms/step  
361/361 [=====] - 7s 20ms/step

The clean accuracy is: 98.53641638520828  
The attack success rate is:  
68%| | 41/60 [11:47<05:00, 15.84s/it]

100.0  
The pruned channel index is: 58  
361/361 [=====] - 9s 24ms/step  
361/361 [=====] - 7s 19ms/step  
70%| | 42/60 [12:05<04:57, 16.54s/it]

The clean accuracy is: 98.19000606218066  
The attack success rate is: 100.0  
The pruned channel index is: 3  
361/361 [=====] - 9s 25ms/step  
361/361 [=====] - 7s 19ms/step  
72%| | 43/60 [12:22<04:43, 16.65s/it]

The clean accuracy is: 97.65307006148784  
The attack success rate is: 100.0  
The pruned channel index is: 42  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 19ms/step  
73%| | 44/60 [12:41<04:35, 17.23s/it]

The clean accuracy is: 97.50584567420108  
The attack success rate is: 100.0  
The pruned channel index is: 1  
361/361 [=====] - 7s 19ms/step

WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet to be built. `model.compile\_metrics` will be empty until you train or evaluate the model.

The accuracy drops at least 2%, saved the model  
361/361 [=====] - 7s 19ms/step  
75%| | 45/60 [12:59<04:22, 17.47s/it]

The clean accuracy is: 95.75647354291158  
The attack success rate is: 100.0  
The pruned channel index is: 29  
361/361 [=====] - 8s 22ms/step  
361/361 [=====] - 7s 19ms/step  
77%| | 46/60 [13:20<04:20, 18.60s/it]

The clean accuracy is: 95.20221702606739  
The attack success rate is: 99.9913397419243  
The pruned channel index is: 16  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 7s 19ms/step

The clean accuracy is: 94.7172425738287  
The attack success rate is: 99.9913397419243  
The pruned channel index is: 56

```

78%|          | 47/60 [13:34<03:45, 17.36s/it]
361/361 [=====] - 7s 19ms/step
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet
to be built. `model.compile_metrics` will be empty until you train or evaluate
the model.

The accuracy drops at least 4%, saved the model
361/361 [=====] - 7s 19ms/step

80%|          | 48/60 [13:49<03:17, 16.49s/it]

The clean accuracy is: 92.09318437689443
The attack success rate is: 99.9913397419243
The pruned channel index is: 46
361/361 [=====] - 7s 19ms/step
361/361 [=====] - 7s 20ms/step

82%|          | 49/60 [14:04<02:57, 16.09s/it]

The clean accuracy is: 91.49562656967177
The attack success rate is: 99.9913397419243
The pruned channel index is: 5
361/361 [=====] - 7s 19ms/step
361/361 [=====] - 8s 21ms/step

83%|          | 50/60 [14:19<02:38, 15.88s/it]

The clean accuracy is: 91.01931237550879
The attack success rate is: 99.98267948384861
The pruned channel index is: 8
361/361 [=====] - 7s 19ms/step
361/361 [=====] - 8s 21ms/step

85%|          | 51/60 [14:35<02:21, 15.72s/it]

The clean accuracy is: 89.17467740538669
The attack success rate is: 80.73958603966398
The pruned channel index is: 11
361/361 [=====] - 7s 19ms/step

WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet
to be built. `model.compile_metrics` will be empty until you train or evaluate
the model.

The accuracy drops at least 10%, saved the model
361/361 [=====] - 7s 19ms/step

87%|          | 52/60 [14:50<02:03, 15.42s/it]

```



The clean accuracy is: 84.43751623798389  
The attack success rate is: 77.015675067117  
The pruned channel index is: 54  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 7s 19ms/step  
88%| | 53/60 [15:04<01:46, 15.24s/it]

The clean accuracy is: 76.48739932449988  
The attack success rate is: 35.71490430414826  
The pruned channel index is: 10  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 7s 19ms/step  
90%| | 54/60 [15:19<01:30, 15.08s/it]

The clean accuracy is: 54.8627349095003  
The attack success rate is: 6.954187234779596  
The pruned channel index is: 28  
361/361 [=====] - 7s 19ms/step  
361/361 [=====] - 7s 19ms/step  
92%| | 55/60 [15:34<01:14, 14.97s/it]

The clean accuracy is: 27.08928726076037  
The attack success rate is: 0.4243526457088421  
The pruned channel index is: 35  
361/361 [=====] - 10s 27ms/step  
361/361 [=====] - 7s 20ms/step  
93%| | 56/60 [15:52<01:03, 15.95s/it]

The clean accuracy is: 13.87373343725643  
The attack success rate is: 0.0  
The pruned channel index is: 18  
361/361 [=====] - 7s 21ms/step  
361/361 [=====] - 7s 20ms/step  
95%| | 57/60 [16:08<00:47, 15.84s/it]

The clean accuracy is: 7.101411622066338  
The attack success rate is: 0.0  
The pruned channel index is: 4  
361/361 [=====] - 7s 20ms/step  
361/361 [=====] - 7s 20ms/step  
97%| | 58/60 [16:23<00:31, 15.75s/it]

```
The clean accuracy is: 0.7188014202823244
The attack success rate is: 0.0
The pruned channel index is: 52
361/361 [=====] - 7s 19ms/step
361/361 [=====] - 7s 20ms/step
100%|          | 60/60 [16:54<00:00, 16.91s/it]
```

**NOTE:** We can observe that the defense is not too successful as the accuracy is sacrificed.


```
clean_accuracy: [98.64899974019225, 98.64899974019225, 98.64899974019225, 98.64899974019225,
98.64899974019225, 98.64899974019225, 98.64899974019225, 98.64899974019225,
98.64899974019225, 98.64899974019225, 98.64899974019225, 98.64899974019225,
98.64899974019225, 98.64899974019225, 98.64899974019225, 98.64899974019225,
98.64899974019225, 98.64899974019225, 98.64899974019225, 98.64899974019225,
98.64899974019225, 98.64899974019225, 98.64899974019225, 98.64899974019225,
98.64899974019225, 98.64899974019225, 98.64899974019225, 98.64899974019225,
98.64899974019225, 98.64899974019225, 98.64899974019225, 98.64899974019225,
98.64899974019225, 98.64899974019225, 98.64033948211657, 98.64033948211657,
98.63167922404088, 98.65765999826795, 98.64899974019225, 98.6056984498138,
98.57105741751104, 98.53641638520828, 98.19000606218066, 97.65307006148784,
97.50584567420108, 95.75647354291158, 95.20221702606739, 94.7172425738287,
92.09318437689443, 91.49562656967177, 91.01931237550879, 89.17467740538669,
84.43751623798389, 76.48739932449988, 54.8627349095003, 27.08928726076037,
13.87373343725643, 7.101411622066338, 1.5501861955486274, 0.7188014202823244,
0.0779423226812159]

attack success rate: [100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0,
100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0,
100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0,
100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0, 100.0,
100.0, 100.0, 100.0, 100.0, 99.9913397419243, 99.9913397419243,
```

```
99.9913397419243, 99.9913397419243, 99.98267948384861, 80.73958603966398,
77.015675067117, 35.71490430414826, 6.954187234779596, 0.4243526457088421, 0.0,
0.0, 0.0, 0.0, 0.0]
```

```
[44]: x_axis = np.arange(1,61)/60
plt.plot(x_axis,clean_acc)
plt.plot(x_axis,asrate)
plt.legend(['clean_accuracy','attack success rate'])
plt.xlabel("fraction of pruned channels")
plt.ylabel("rate")
plt.title("accuracy and attack success rate for validation dataset")
```

```
[44]: Text(0.5, 1.0, 'accuracy and attack success rate for validation dataset')
```



```
[45]: index = np.where(np.array(clean_acc) <= (clean_data_acc-30))[0]
print("The attack success rate when the accuracy drops at least 30%: ",
      ↪asrate[index[0]])
```

The attack success rate when the accuracy drops at least 30%: 6.954187234779596

## 1.4 Combining the models

Here we combine two models which are  $B$  (original badnet model) and  $B'$  (pruned model). The combined model is the *goodnet*. If the predictions from  $B$  and  $B'$  are the same then the *goodnet* will output the prediction.

```
[46]: class G(keras.Model):
    def __init__(self, B, B_prime):
        super(G, self).__init__()
        self.B = B
        self.B_prime = B_prime

    def predict(self, data):
        y = np.argmax(self.B(data), axis=1)
        y_prime = np.argmax(self.B_prime(data), axis=1)
        pred = np.zeros(data.shape[0])
        for i in range(data.shape[0]):
            if y[i]==y_prime[i]:
                pred[i] = y[i]
            else:
                pred[i] = 1283
        return pred
```

## 1.5 Evaluate the combined model

```
[48]: test_data_filename = '/content/drive/MyDrive/Lab2_mm11070/lab3/data/c1/test.h5'
poisoned_test_data_filename = '/content/drive/MyDrive/Lab2_mm11070/lab3/data/bd/
↳bd_test.h5'
test_model_X_2_filename = '/content/model_X=2.h5'
test_model_X_4_filename = '/content/model_X=4.h5'
test_model_X_10_filename = '/content/model_X=10.h5'
```

```
[49]: test_model_X_2 = keras.models.load_model(test_model_X_2_filename)
test_model_X_4 = keras.models.load_model(test_model_X_4_filename)
test_model_X_10 = keras.models.load_model(test_model_X_10_filename)
```

WARNING:tensorflow:No training configuration found in the save file, so the model was \*not\* compiled. Compile it manually.

WARNING:tensorflow:No training configuration found in the save file, so the model was \*not\* compiled. Compile it manually.

WARNING:tensorflow:No training configuration found in the save file, so the model was \*not\* compiled. Compile it manually.

Ignore the warnings

```
[50]: x_test_data, y_test_data = data_loader(test_data_filename)
x_test_poisoned_data, y_test_poisoned_data =
↳data_loader(poisoned_test_data_filename)
```

```
[51]: print("x_test_data shape: ",x_test_data.shape)
print("x_test_poisoned data shape: ",x_test_poisoned_data.shape)
```

x\_test\_data shape: (12830, 55, 47, 3)

x\_test\_poisoned data shape: (12830, 55, 47, 3)

```
[52]: G_model_X_2 = G(model, test_model_X_2)
      G_model_X_4 = G(model, test_model_X_4)
      G_model_X_10 = G(model, test_model_X_10)
```

### 1.5.1 Evaluating on the test dataset

```
[53]: cl_test_2_label_p = np.argmax(test_model_X_2.predict(x_test_data), axis=1)
      clean_test_2_accuracy = np.mean(np.equal(cl_test_2_label_p, y_test_data))*100
      print('2% drops model, the clean test data Classification accuracy:',
            ↪clean_test_2_accuracy)

      bd_test_2_label_p = np.argmax(test_model_X_2.predict(x_test_poisoned_data),
            ↪axis=1)
      asr_2 = np.mean(np.equal(bd_test_2_label_p, y_test_poisoned_data))*100
      print('2% drops model, Attack Success Rate:', asr_2)

      cl_test_4_label_p = np.argmax(test_model_X_4.predict(x_test_data), axis=1)
      clean_test_4_accuracy = np.mean(np.equal(cl_test_4_label_p, y_test_data))*100
      print('4% drops model, the clean test data classification accuracy:',
            ↪clean_test_4_accuracy)

      bd_test_4_label_p = np.argmax(test_model_X_4.predict(x_test_poisoned_data),
            ↪axis=1)
      asr_4 = np.mean(np.equal(bd_test_4_label_p, y_test_poisoned_data))*100
      print('4% drops model, Attack Success Rate:', asr_4)

      cl_test_10_label_p = np.argmax(test_model_X_10.predict(x_test_data), axis=1)
      clean_test_10_accuracy = np.mean(np.equal(cl_test_10_label_p, y_test_data))*100
      print('10% drops model, the clean test data classification accuracy:',
            ↪clean_test_10_accuracy)

      bd_test_10_label_p = np.argmax(test_model_X_10.predict(x_test_poisoned_data),
            ↪axis=1)
      asr_10 = np.mean(np.equal(bd_test_10_label_p, y_test_poisoned_data))*100
      print('10% drops model, Attack Success Rate:', asr_10)
```

```
401/401 [=====] - 8s 20ms/step
2% drops model, the clean test data Classification accuracy: 95.90023382696803
401/401 [=====] - 8s 19ms/step
2% drops model, Attack Success Rate: 100.0
401/401 [=====] - 8s 19ms/step
4% drops model, the clean test data classification accuracy: 92.29150428682775
401/401 [=====] - 8s 19ms/step
4% drops model, Attack Success Rate: 99.98441153546376
401/401 [=====] - 8s 19ms/step
10% drops model, the clean test data classification accuracy: 84.54403741231489
401/401 [=====] - 8s 19ms/step
```

10% drops model, Attack Success Rate: 77.20966484801247

## 1.5.2 Summarizing the fixed models

Accuracy vs Attack Rate

```
[54]: test_acc = [clean_test_2_accuracy, clean_test_4_accuracy,
    ↪ clean_test_10_accuracy]
attack_rate = [asr_2, asr_4, asr_10]
data = {
    "text_acc": test_acc,
    "attack_rate": attack_rate,
    "model": ["repaired_2%", "repaired_4%", "repaired_10%"]
}
df = pd.DataFrame(data)
df.set_index('model')
```

```
[54]:
```

	text_acc	attack_rate
model		
repaired_2%	95.900234	100.000000
repaired_4%	92.291504	99.984412
repaired_10%	84.544037	77.209665


```
[65]: opacity = 0.5
bar_width = 0.38

plt.xlabel('% drops model')
plt.ylabel('Rate')

plt.xticks(range(len(test_acc)), ('2%', '4%', '10%'))
bar1 = plt.bar(np.arange(len(test_acc)) + bar_width, test_acc, bar_width,
    ↪ align='center', alpha=opacity, color='g', label='accuracy')
bar2 = plt.bar(range(len(attack_rate)), attack_rate, bar_width, align='center',
    ↪ alpha=opacity, color='b', label='attack rate')

# Adding value above bar
for rect in bar1 + bar2:
    height = rect.get_height()
    plt.text(rect.get_x() + rect.get_width() / 2.0, height, f'{height:.02f}',
    ↪ ha='center', va='bottom')

plt.legend(bbox_to_anchor=(1.4, 1))
plt.tight_layout()
plt.title('Performance of Repaired Model')
sns.despine()
plt.show()
```



These are the *goodnets* which combines the two models that are the original badnet and the ‘fixed’ model

```
[56]: G_cl_test_2_label_p = G_model_X_2.predict(x_test_data)
G_clean_test_2_accuracy = np.mean(np.equal(cl_test_2_label_p, y_test_data))*100
print('Combined 2% drops model, the clean test data Classification accuracy:',
      ↪G_clean_test_2_accuracy)

G_bd_test_2_label_p = G_model_X_2.predict(x_test_poisoned_data)
G_asr_2 = np.mean(np.equal(bd_test_2_label_p, y_test_poisoned_data))*100
print('Combined 2% drops model, Attack Success Rate:', G_asr_2)

G_cl_test_4_label_p = G_model_X_4.predict(x_test_data)
G_clean_test_4_accuracy = np.mean(np.equal(cl_test_4_label_p, y_test_data))*100
print('Combined 4% drops model, the clean test data Classification accuracy:',
      ↪G_clean_test_4_accuracy)

G_bd_test_4_label_p = G_model_X_4.predict(x_test_poisoned_data)
G_asr_4 = np.mean(np.equal(bd_test_4_label_p, y_test_poisoned_data))*100
print('Combined 4% drops model, Attack Success Rate:', G_asr_4)

G_cl_test_10_label_p = G_model_X_10.predict(x_test_data)
```

```

G_clean_test_10_accuracy = np.mean(np.equal(cl_test_10_label_p,
    ↪y_test_data))*100
print('Combined 10% drops model, the clean test data Classification accuracy:',
    ↪G_clean_test_10_accuracy)

G_bd_test_10_label_p = G_model_X_10.predict(x_test_poisoned_data)
G_asr_10 = np.mean(np.equal(bd_test_10_label_p, y_test_poisoned_data))*100
print('Combined 10% drops model, Attack Success Rate:', G_asr_10)

```

Combined 2% drops model, the clean test data Classification accuracy:  
 95.90023382696803  
 Combined 2% drops model, Attack Success Rate: 100.0  
 Combined 4% drops model, the clean test data Classification accuracy:  
 92.29150428682775  
 Combined 4% drops model, Attack Success Rate: 99.98441153546376  
 Combined 10% drops model, the clean test data Classification accuracy:  
 84.54403741231489  
 Combined 10% drops model, Attack Success Rate: 77.20966484801247

```

[57]: G_test_acc = [G_clean_test_2_accuracy, G_clean_test_4_accuracy,
    ↪G_clean_test_10_accuracy]
G_attack_rate = [G_asr_2, G_asr_4, G_asr_10]
G_data = {
    "G_text_acc": G_test_acc,
    "G_attack_rate": G_attack_rate,
    "G_model": ["G_2%", "G_4%", "G_10%"]
}
G_df = pd.DataFrame(G_data)
G_df.set_index('G_model')

```

```

[57]:      G_text_acc  G_attack_rate
G_model
G_2%      95.900234      100.000000
G_4%      92.291504      99.984412
G_10%     84.544037      77.209665

```

```

[66]: opacity = 0.5
      bar_width = 0.38

      plt.xlabel('combined % drops model')
      plt.ylabel('Rate')

      plt.xticks(range(len(G_test_acc)), ('2%', '4%', '10%'))
      bar1 = plt.bar(np.arange(len(G_test_acc)) + bar_width, G_test_acc, bar_width,
    ↪align='center', alpha=opacity, color='g', label='accuracy')
      bar2 = plt.bar(range(len(G_attack_rate)), G_attack_rate, bar_width,
    ↪align='center', alpha=opacity, color='b', label='attack rate')

```




```

for rect in bar1 + bar2:
    height = rect.get_height()
    plt.text(rect.get_x() + rect.get_width() / 2.0, height, f'{height:.02f}',
             ha='center', va='bottom')

plt.legend(bbox_to_anchor=(1.4, 1))
plt.tight_layout()
plt.title('Performance of GoodNet Model')
sns.despine()
plt.show()

```



[ ]: